

2.03. knn

October 14, 2024

1 Predictive Modeling

```
[17]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import httpimport
import joblib

from imblearn.over_sampling import SMOTE
from pathlib import Path
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.model_selection import train_test_split, RandomizedSearchCV,
↳ GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
```

```
[18]: # Import personal library
with httpimport.github_repo("junclemente", "jcds", ref="master"):
    import jcds.metrics as jm
```

```
[19]: # Import datasets
datasets = Path("../datasets")
train_data = "training_data.csv"
val_data = "validation_data.csv"
test_data = "testing_data.csv"
train_df = pd.read_csv(datasets / train_data)
val_df = pd.read_csv(datasets / val_data)
test_df = pd.read_csv(datasets / test_data)
display(train_df.head())
display(val_df.head())
display(test_df.head())
```

	Undergrad_Degree	Work_Experience	Employability_Before	Status	\
0	Computer Science	No	185.174286	Placed	
1	Engineering	No	206.867959	Not Placed	
2	Art	No	234.881837	Not Placed	

3	Finance	No	173.900408	Placed
4	Art	No	184.063980	Not Placed

	Status_enc
0	1
1	0
2	0
3	1
4	0

	Undergrad_Degree	Work_Experience	Employability_Before	Status \
0	Business	Yes	261.272959	Placed
1	Engineering	No	173.558776	Not Placed
2	Finance	No	205.074388	Placed
3	Business	Yes	230.526020	Placed
4	Business	No	229.000000	Not Placed

	Status_enc
0	1
1	0
2	1
3	1
4	0

	Undergrad_Degree	Work_Experience	Employability_Before	Status \
0	Finance	No	168.775918	Placed
1	Business	Yes	195.508673	Placed
2	Computer Science	No	260.760510	Placed
3	Art	No	231.892551	Not Placed
4	Computer Science	Yes	400.000000	Placed

	Status_enc
0	1
1	1
2	1
3	0
4	1

1.1 Setup Training and Validation dataframes

```
[20]: # Variables to use for predictive modeling
variables = ["Undergrad_Degree", "Work_Experience", "Employability_Before"]
target = "Status_enc"
```

```
[21]: # Setup train, val, and test dataframes
X_train = train_df[variables]
y_train = train_df[target]
```

```

X_val = val_df[variables]
y_val = val_df[target]

X_test = test_df[variables]
y_test = test_df[target]

# One-hot encode categorical variables
X_train = pd.get_dummies(X_train, drop_first=True)
X_val = pd.get_dummies(X_val, drop_first=True)
X_test = pd.get_dummies(X_test, drop_first=True)

# Standardize cont / Initialize scaler
scaler = StandardScaler()
std_cols = ["Employability_Before"]
X_train[std_cols] = scaler.fit_transform(X_train[std_cols])
X_val[std_cols] = scaler.transform(X_val[std_cols])
X_test[std_cols] = scaler.transform(X_test[std_cols])

display(X_train.head())
display(X_val.head())
display(X_test.head())

```

	Employability_Before	Undergrad_Degree_Business \
0	-0.800730	False
1	-0.244034	False
2	0.474848	False
3	-1.090036	False
4	-0.829222	False

	Undergrad_Degree_Computer Science	Undergrad_Degree_Engineering \
0	True	False
1	False	True
2	False	False
3	False	False
4	False	False

	Undergrad_Degree_Finance	Work_Experience_Yes
0	False	False
1	False	False
2	False	False
3	True	False
4	False	False

	Employability_Before	Undergrad_Degree_Business \
0	1.152088	True
1	-1.098803	False
2	-0.290060	False
3	0.363071	True

4	0.323911	True
---	----------	------

	Undergrad_Degree_Computer Science	Undergrad_Degree_Engineering \
0	False	False
1	False	True
2	False	False
3	False	False
4	False	False

	Undergrad_Degree_Finance	Work_Experience_Yes
0	False	True
1	False	False
2	True	False
3	False	True
4	False	False

	Employability_Before	Undergrad_Degree_Business \
0	-1.221539	False
1	-0.535532	True
2	1.138938	False
3	0.398138	False
4	4.712054	False

	Undergrad_Degree_Computer Science	Undergrad_Degree_Engineering \
0	False	False
1	False	False
2	True	False
3	False	False
4	True	False

	Undergrad_Degree_Finance	Work_Experience_Yes
0	True	False
1	False	True
2	False	False
3	False	False
4	False	True

1.1.1 Check class balance

```
[22]: y_train.value_counts()
```

```
[22]: Status_enc
1    348
0    228
Name: count, dtype: int64
```

```
[23]: # Use SMOTE to balance classes
smote = SMOTE(random_state=42)
```

```
X_train, y_train = smote.fit_resample(X_train, y_train)

y_train.value_counts()
```

```
[23]: Status_enc
1      348
0      348
Name: count, dtype: int64
```

2 K-Nearest Neighbor

2.1 RandomSearchCV

```
[24]: knn = KNeighborsClassifier()
param_dist = {
    "n_neighbors": [2, 4, 6, 8, 10, 12],
    "weights": ["uniform", "distance"],
    "metric": ["euclidean", "manhattan", "chebyshev"],
}

random_search = RandomizedSearchCV(
    estimator=knn,
    param_distributions=param_dist,
    n_iter=20,
    cv=5,
    scoring="accuracy",
    random_state=42,
)

random_search.fit(X_train, y_train)

print(random_search.best_params_)

{'weights': 'uniform', 'n_neighbors': 4, 'metric': 'chebyshev'}
```

2.2 GridSearchCV

```
[25]: param_grid = {
    "n_neighbors": [
        2,
        3,
        4,
        5,
        6,
        7,
    ],
    "weights": ["uniform", "distance"],
    "metric": ["euclidean", "manhattan", "chebyshev"],
}
```

```

}

grid_search = GridSearchCV(
    estimator=knn, param_grid=param_grid, cv=5, scoring="accuracy"
)

grid_search.fit(X_train, y_train)
print(grid_search.best_params_)

{'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'uniform'}

```

2.3 Prediction Model

```

[26]: knn_model = KNeighborsClassifier(
        metric="manhattan",
        n_neighbors=3,
        weights="uniform",
    )
knn_model.fit(X_train, y_train)

y_pred = knn_model.predict(X_val)

```

```

[27]: cm = confusion_matrix(y_val, y_pred)
jm.mc_confusion(cm)

```

Confusion Matrix:

```

[[153   8]
 [  6 217]]

```

```

[27]:

```

	Class 0	Class 1
Accuracy	0.96354	0.96354
Error rate	0.03646	0.03646
Sensitivity (Recall)	0.95031	0.97309
Specificity	0.97309	0.95031
Precision	0.96226	0.96444
F1	0.95625	0.96875
F2	0.95268	0.97135
F0.5	0.95985	0.96616

2.4 Test

```

[28]: test_pred = knn_model.predict(X_test)
cm_test = confusion_matrix(y_test, test_pred)
jm.mc_confusion(cm_test)

```

Confusion Matrix:

```

[[ 92   3]
 [  5 140]]

```

```
[28]:
```

	Class 0	Class 1
Accuracy	0.96667	0.96667
Error rate	0.03333	0.03333
Sensitivity (Recall)	0.96842	0.96552
Specificity	0.96552	0.96842
Precision	0.94845	0.97902
F1	0.95833	0.97222
F2	0.96436	0.96819
F0.5	0.95238	0.97629

3 Export Model

```
[29]: models = Path("../models")  
      joblib.dump(knn_model, models / "k_nearest_neighbor_model.pkl")
```

```
[29]: ['../models/k_nearest_neighbor_model.pkl']
```

```
[ ]:
```